

Indexing Large Data Sets with Watsonx Discovery S3 Connector

Introduction

Watsonx Discovery comes with a lot of capabilities providing a lot of features and connectors to leverage from. While the Q&A with RAG Accelerator provides an efficient approach for data ingestion and indexing, indexing of very large datasets (terabytes and so) can be enhanced using connector feature provided by Watsonx Discovery. Large datasets can be stored in Amazon S3 bucket and using an S3 connector with Watsonx Discovery, documents can be directly ingested and indexed to Elasticsearch vector store.

Key tasks involved in Indexing Large Datasets from S3:

- **Setup self-managed S3 connector** to connect your Amazon S3 bucket to Watsonx Discovery.
- **Setup Elasticsearch pipeline** for chunking and indexing of the documents.
- **Setup self-managed Content Extraction Service** to process large documents (e.g., documents exceeding 100MB).

Prerequisites

- The datasets to be indexed are present in AWS S3 bucket in pdf format.
- An AWS EC2 instance is available to execute the S3 connector.

Processing Compressed Files with AWS Lambda and S3 Connector

Since the S3 connector cannot directly process compressed files, you will first need to extract the files from the compressed archives and store them in a new S3 bucket. AWS Lambda can handle this extraction automatically when a new compressed file (such as .tar or .tar.gz) is uploaded to the source S3 bucket.

Steps Overview:

- Create AWS IAM Roles and Policies for Lambda to access and interact with S3.
- Create a Lambda function to extract and store the files from compressed archives. Here is a sample Lambda function written in Python that handles the extraction of .tar files from an S3 bucket:

```
import boto3
import botocore
import tarfile

from io import BytesIO
s3_client = boto3.client('s3')

def lambda_handler(event, context):

    bucket = event['Records'][0]['s3']['bucket']['name']
    key = event['Records'][0]['s3']['object']['key']
    destination_bucket = bucket+'_extracted'
```

```

        input_tar_file = s3_client.get_object(Bucket = bucket, Key =
key)
        input_tar_content = input_tar_file['Body'].read()

        dest_key = key.split('.')[0]+'/'

        with tarfile.open(fileobj = BytesIO(input_tar_content)) as
tar:
            for tar_resource in tar:
                if (tar_resource.isfile()):
                    inner_file_bytes =
tar.extractfile(tar_resource).read()
                    print('extracted file: ',tar_resource.name)
                    tar_resource_file = tar_resource.name.split('/')
[1]
                    print('tar_resource_file: ', tar_resource_file)

s3_client.upload_fileobj(BytesIO(inner_file_bytes),
                           Bucket = destination_bucket,
                           Key = dest_key+tar_resource_file)
print('file uploaded')

```

- Set up the trigger for Lambda to automatically execute upon file upload in the S3 bucket.
- Create a destination S3 bucket to store the extracted files.

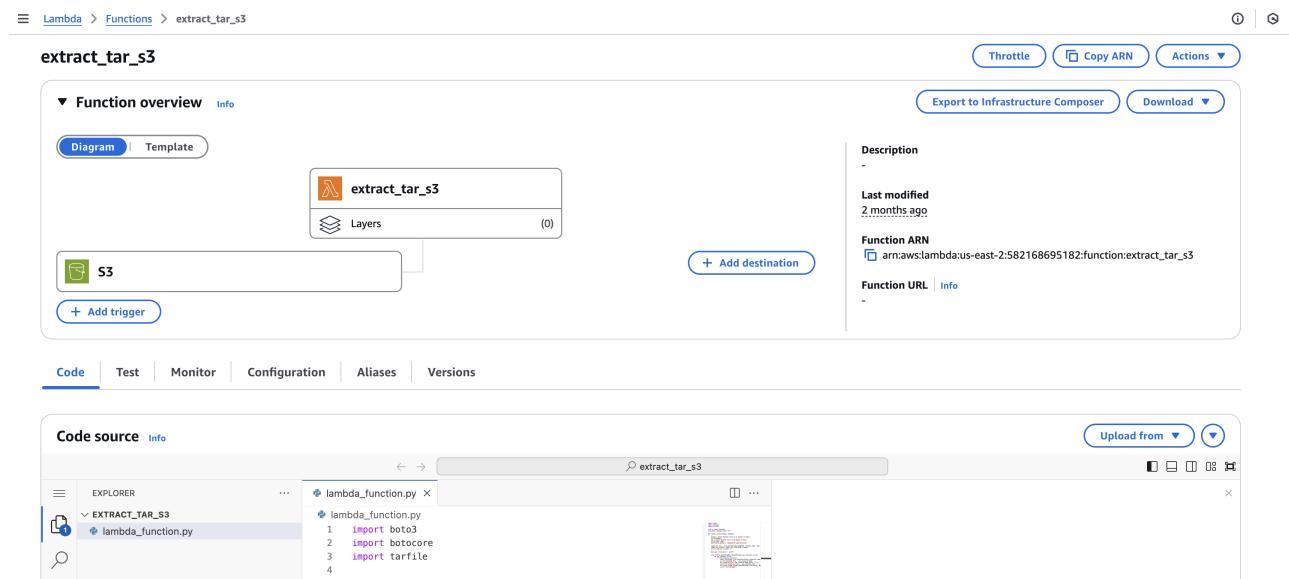


Figure: AWS Lambda Function

- Upon execution of the above created Lambda function, the destination bucket will get populated with the PDF files extracted from the .tar archive, preserving the same directory structure as the source S3 bucket.

Setting up the S3 Connector on wx-discovery Kibana

To integrate the pipeline, the S3 connector needs to be set up within wx-discovery Kibana. Follow the steps below to configure the S3 connector:

1. Navigate to Kibana:

- o Access wx-discovery Kibana and go to Elasticsearch.
- o In the left-hand menu, click on Connectors.

2. Create a New Connector:

- o Click on New Connector to begin the setup process.
- o Select S3 as the connector source.

3. Configure the Connector:

- o A default connector name will be auto-generated, which you can update as needed.
- o Choose Setup as self-managed.

The screenshot shows the Kibana interface for creating a new connector. The top navigation bar includes 'Elasticsearch', 'Content', 'Connectors', and 'New connector'. On the left, a sidebar lists various features like Home, Content, Indices, Connectors (which is selected), Web crawlers, Build, and Relevance. The main area is titled 'Create a connector' with the sub-instruction 'Extract, transform, index and sync data from a third-party data source.' Below this is a vertical navigation bar with steps: 'Start' (selected), 'Deployment', 'Configuration', and 'Finish up'. A note says 'Elastic S3 connector reference'. To the right of this is the 'Start' configuration panel. It shows 'Connector' set to 'S3' and 'Connector name' set to 's3-ac1d' (with a note that it should be lowercase and no spaces). There's a 'Description' field and an 'Optional' link. Below this is the 'Setup' section, which asks 'Choose how to deploy and manage your connector:' with 'Self-managed' selected. The 'Deployment' section notes that configurations will be automatically handled. At the bottom is a blue 'Next' button.

Figure: Kibana Connector Setup Step

- Click Deployment and then select Next to continue.
- Choose "Run with Docker" as the deployment method.

4. Connector ID:

- o A new connector_id will be generated during the setup.

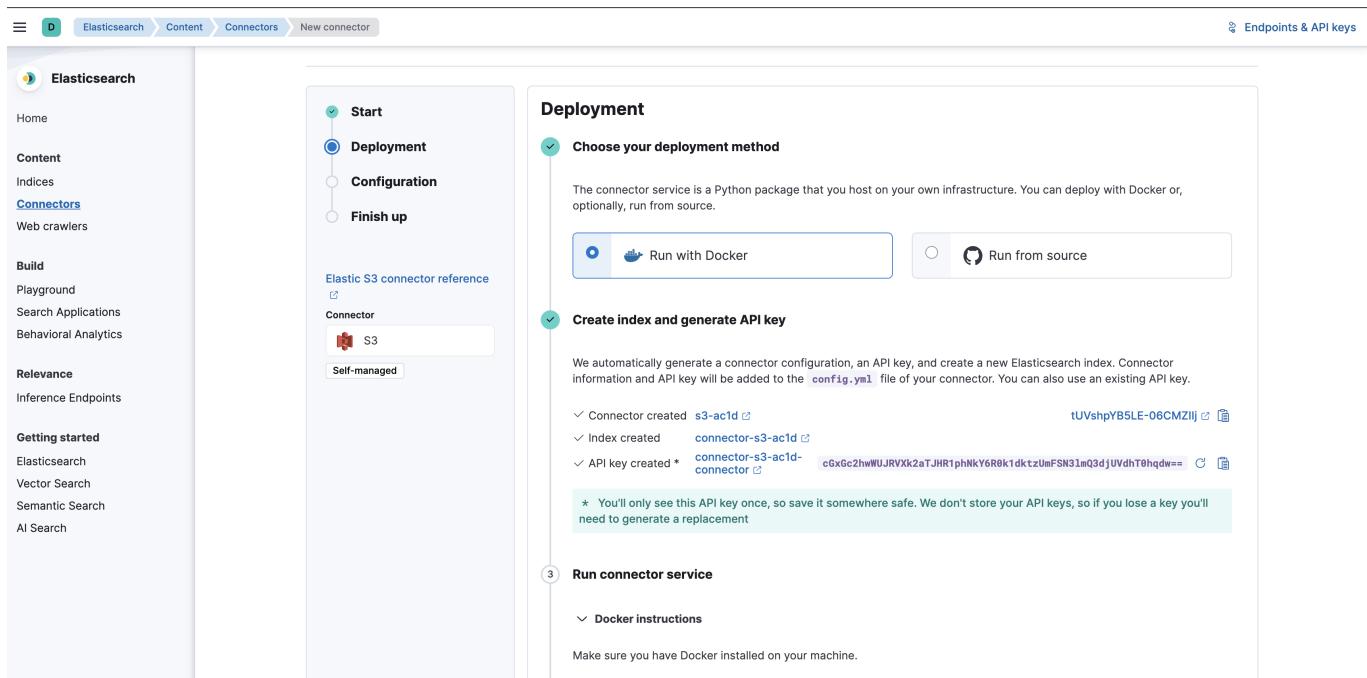


Figure: Kibana Connector Setup Step

- Copy this connector_id and will be used in the config.yaml file for further configuration.

5. Start the Connector:

- o Once the setup is complete, start the connector from the AWS Console.
 - To set up and start the connector on the AWS Console, please follow the steps as outlined in the section **Steps for Setting Up the S3 Connector Instance** below.
- o The connector status should now appear as Connected in the AWS Console.

Connector name	Index name	Docs count	Connector type	Ingestion status	Actions
s3-bdaa	connector-s3-bdaa	161	S3	Connected	

Figure: Kibana Connector Setup Step

6. Configure the connector

- o Go to the S3 Connector configuration page.
- o Navigate to Configuration.
- o Click Edit Configuration.

- Add the AWS S3 bucket name containing the input PDF files, along with the AWS Access Key and AWS Secret Key.
- Save the configuration.

Figure: Kibana Connector Setup Configuration

By following these steps, the S3 connector will be properly configured and ready for integration with the pipeline in wx-discovery Kibana.

Setting Up Elasticsearch Pipeline for Chunking and Indexing Documents

To set up the Elasticsearch pipeline for chunking and indexing the documents, follow these steps:

1. Create a New Elasticsearch Index:

First, create a new Elasticsearch index with sharding enabled and an increased field limit. Optimal number of shards can be derived based on the input data size and cluster configuration. Following is the sample settings:

```
{
  "settings": {
    "number_of_shards": 7,
    "index": {
      "refresh_interval": "30s",
      "mapping.total_fields.limit": 5000000
    }
  }
}
```

2. Add the New Index to the S3 Connector:

After creating the new index, add the shard_index to the S3 connector in Kibana. This will enable the connector to push the documents to the newly created index.

3. Add Elasticsearch Pipelines for Nested Chunking and Indexing:

To ensure that the documents are chunked and indexed properly, add the following Elasticsearch

pipelines to the connector, based on the instructions from the [Watson Developer Cloud Assistant Toolkit](#).

The screenshot shows the 's3-caed' pipeline configuration page. On the left sidebar, under 'Content' > 'Connectors', 's3-caed' is selected. The main area has tabs for Overview, Documents, Mappings, Scheduling, Pipelines (selected), and Configuration. Under 'Ingest Pipelines', there are two entries: 'connector-s3-caed' (Managed) and 'connector-s3-caed@custom' (Edit pipeline, 0 Processors). A note says 'Custom ingest pipeline for connector-s3-caed'. Below this is a section for 'Machine Learning Inference Pipelines' with a note: 'Inference pipelines will be run as processors from the Search Ingest Pipeline'. A green message at the bottom says 'Your ELSER model has started.'.

Figure: Elasticsearch pipeline

4. Update the Script Pipeline Processor:

Update the Script Pipeline Processor with the following source code to perform the chunking of documents, handle overlapping passages and handle empty body content:

```
if(ctx['body'] != null)
{
    String[] envSplit = /((?<!M(r|s|rs)\.)(?=<\.)|(?=<\!)|(?=<\?))/
                    .split(ctx['body']);
    ctx['passages'] = [];

    StringBuilder overlappingText = new StringBuilder();

    int title_subs_idx = 200;
    if(ctx['body'].splitOnToken('_')[0].length() < 200)
        title_subs_idx = ctx['body'].splitOnToken('_')[0].length();
    int i = 0;
    while (i < envSplit.length) {
        StringBuilder passageText = new StringBuilder(envSplit[i]);
        int accumLength = envSplit[i].length();
        ArrayList accumLengths = [];
        accumLengths.add(accumLength);

        int j = i + 1;
        while (j < envSplit.length && passageText.length() +
envSplit[j].length() < params.model_limit) {
            passageText.append(' ').append(envSplit[j]);
            accumLength += envSplit[j].length();
            accumLengths.add(accumLength);
            j++;
        }
    }
}
```

```

        ctx['passages'].add(['text': overlappingText.toString() +
passageText.toString(), 'title': ctx['body'].splitOnToken('_')
[0].substring(0, title_subs_idx), 'url': ctx['filename']]);
        def startLength = passageText.length() * (1 -
params.overlap_percentage) + 1;

        int k = Collections.binarySearch(accumLengths,
(int)startLength);
        if (k < 0) {
            k = -k - 1;
        }
        overlappingText = new StringBuilder();
        for (int l = i + k; l < j; l++) {
            overlappingText.append(envSplit[l]).append(' ');
        }

        i = j;
    }
}
else{
    ctx['passages'] = [];
}

```

5. Enable Bulk Request Processing

To optimize document indexing, enable bulk request processing in the pipeline. This allows Elasticsearch to process multiple document chunks in a single request, improving performance.

6. Pipeline Completion

After the pipeline setup is complete, the S3 connector will chunk and index the documents into Elasticsearch as per the defined settings.

Steps for Setting Up the S3 Connector Instance

This guide outlines the process for setting up the S3 connector instance without Content Extraction.

Prerequisites:

- AWS Console Access: You need access to the AWS EC2 console to launch an instance.
- Key Pair: Ensure you have the key pair for SSH access.
- Security Group: SSH access is required.

NOTE: All commands provided are based on Ubuntu OS. If using a different operating system, the commands may need to be updated accordingly.

Steps for Setup:

1. Launch an Instance on AWS EC2

- Go to AWS EC2 service.
- Launch a new instance based on your connector requirements (this will depend on CPU, memory, and the number of instances you need).
- Select the Application and OS Image - Ubuntu

- Select the Instance Type (based on connector requirements).
- Choose the key pair.
- Select the appropriate security group (for SSH access).

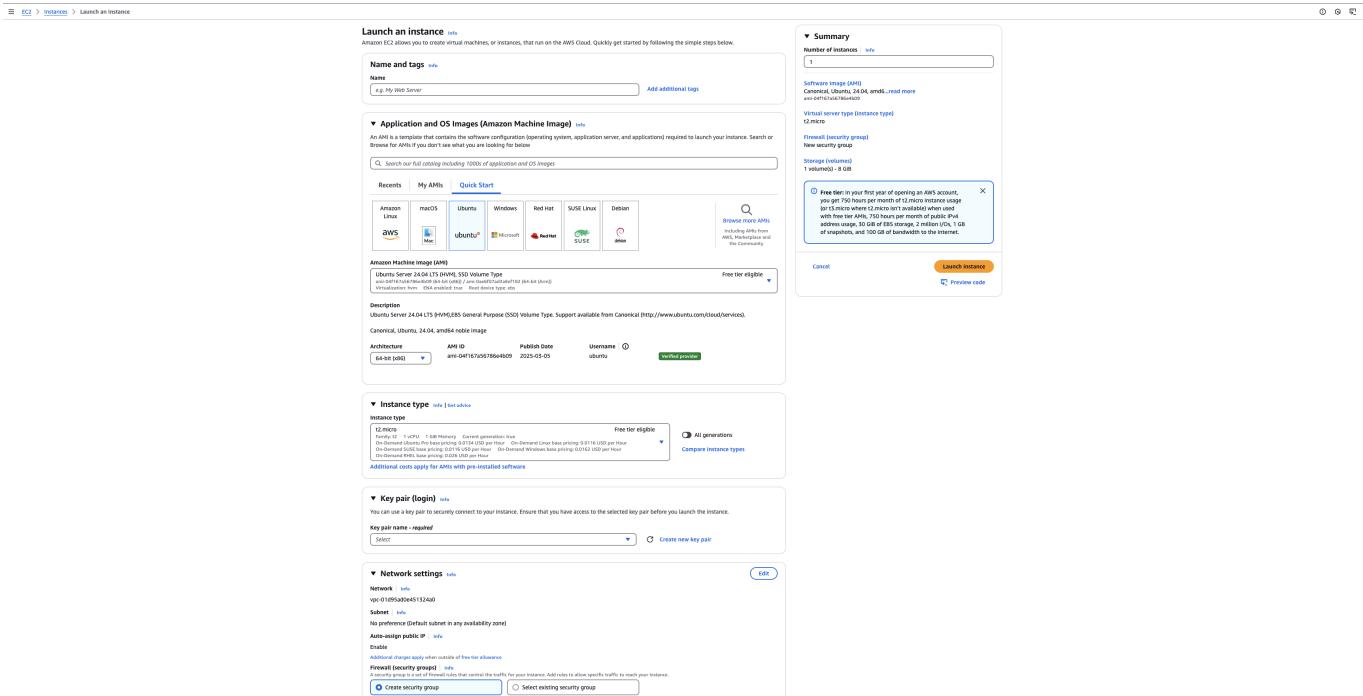


Figure: AWS EC2 instance

2. Access the Instance

- Once the instance is up and running, SSH into the instance.

3. Inside the Container: Install Prerequisites

- Update the system and install required tools:

```
sudo apt install git
sudo apt-get install make
```

4. Clone the Connector Repository

- Clone the connector repository from GitHub:

```
git clone https://github.com/elastic/connectors
cd connectors && git checkout v8.17.0
```

5. Create the config.yml File

- Inside the connectors directory, create the config.yml file:

```
touch config.yml
```

- Add the following configuration details to the config.yml file:

```
connectors:  
  - connector_id: ""  
    service_type: "s3"  
elasticsearch:  
  host: ""  
  username: ""  
  password: ""  
  verify_certs: true  
  ca_certs: /home/ubuntu/connectors/cert.crt
```

6. Create the cert.crt File

- You need to create a certificate file (cert.crt) as specified in the configuration: Ensure the certificate file is available at the path /home/ubuntu/connectors/cert.crt.

7. Install Python and Prerequisites

- Before starting the connector, install the necessary prerequisites:

```
sudo apt update  
sudo apt upgrade  
sudo apt install python3.12-venv
```

8. Start the Connector

- To start the connector, run the following commands:

```
make install  
make run
```

9. Verify the Connector in Kibana (wx-discovery)

- Go to Kibana under the wx-discovery environment.
- The connector should now show as "connected".

Important Note: File Size Limitation for Above Connector

By default, the Above connector is designed to handle a maximum file size of 10MB. If you attempt to process a file larger than this size, you will encounter issues.

What to Do If Your File Exceeds the Limit

If your input file size exceeds the 10MB limit, you will need to make additional updates to the S3 connector code or configuration to accommodate larger files. This will ensure that your files are processed without errors.

Updating S3 Connector to Process Files Up to 100MB

Follow these steps to update the connector to handle files up to 100MB:

1. Update the DEFAULT_MAX_FILE_SIZE

- o Navigate to the connectors directory:

```
cd connectors
```

- o Open the config.py file for editing:

```
vi connectors/config.py
```

- o Update the DEFAULT_MAX_FILE_SIZE to 100MB (in bytes: 104857600):

```
DEFAULT_MAX_FILE_SIZE = 104857600
```

2. Execute the Connector for Document Ingestion in the Elasticsearch Pipeline

Run the following commands to install and start the connector after making the configuration changes:

```
make install  
make run
```

Important Note:

For documents exceeding 100MB, it cannot be processed directly by the connector as HTTP has a max file size limit of 100MB. To process these larger files, you must set up and integrate a self-managed Content Extraction Service.

Setting Up Content Extraction for S3 Connector

To handle documents larger than 100MB, you must set up a self-managed Content Extraction Service.

Steps to Set Up the Content Extraction Service:

1. Install Docker

Ensure that Docker is installed on the system running the S3 connector. If Docker is not installed, you can install it using the following command:

```
sudo apt install docker.io
```

2. Start the Data Extraction Service in Docker

Start the Data Extraction Service Docker container, which will run on port 8090. The container will allow you to extract content from documents larger than 100MB.

Run the following command to start the extraction service:

```
sudo docker run -d -p 8090:8090 -it -v /app/files:/app/files --name extraction-service docker.elastic.co/integrations/data-extraction-service:0.3.4
```

This command pulls and runs the data-extraction-service Docker image, maps the necessary directories, and exposes port 8090 for the service.

3. Update the S3 Connector Configuration

Now, update the S3 Connector configuration to integrate with the Data Extraction Service.

- Navigate to the connectors directory:

```
cd connectors
```

- Open the config.py file for editing:

```
vi connectors/config.py
```

- Add the Content Extraction Service configuration block: Add the following configuration for the extraction_service to connect the S3 connector to the extraction service:

```
# data-extraction-service settings
extraction_service:
    host: http://localhost:8090
    timeout: 900
    use_file_pointers: true
    shared_volume_dir: '/app/files'
```

4. Enable Content Extraction in the S3 Connector

Next, enable the content extraction feature in the S3 connector's configuration file.

- Navigate to the sources/s3.py file in the connector directory:

```
cd connector/sources
vi s3.py
```

- Update the default configuration to enable the text extraction service by setting `use_text_extraction_service` to True:

```
use_text_extraction_service = True
```

5. Execute the Connector for Document Ingestion in Elasticsearch Pipeline

After making the necessary updates to the configuration, run the following commands to install and start the S3 connector for document ingestion:

```
make install
make run
```

Enabling Content Extraction in Kibana

After making the necessary configuration changes to the S3 connector, you'll need to enable content extraction in Kibana for the connector to properly process documents.

1. In Kibana UI

- Go to the S3 Connector configuration page.
- Navigate to Configuration.
- Click Edit Configuration.
- Enable the toggle for "Use Text Extraction Service".
- Save the configuration.

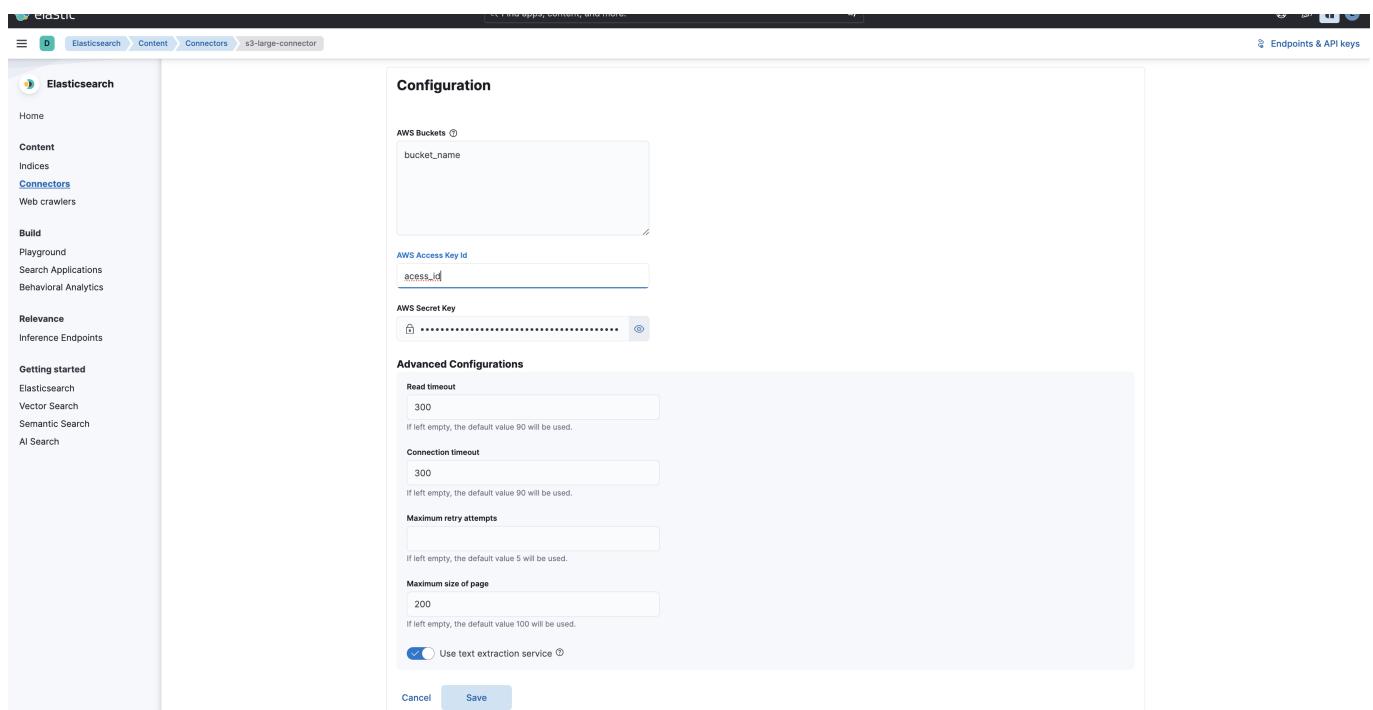


Figure: Kibana Connector Setup Configuration

2. Start the Connector in Kibana

Once content extraction is enabled in Kibana, start the connector to begin processing documents with content extraction enabled.

License

Sample Materials, provided under [license](#).

Licensed Materials - Property of IBM.

© Copyright IBM Corp. 2024, 2025. All Rights Reserved.

US Government Users Restricted Rights - Use, duplication or disclosure restricted by GSA ADP

Schedule Contract with IBM Corp.