

3.3 RECAP & NEXT STEPS (2H DISCUSSION)

SESSION OVERVIEW

This is a 2-hour interactive recap and discussion to consolidate learning and plan next steps.

PART 1 – TIMELINE RECAP

Walk through the workshop story:

- **Day 0:** Environment setup
 - Two env repos (Ollama + watsonx).
 - accelerator repo.
 - labs-src reference notebooks.
- **Day 1:** LLM basics & prompting
 - Prompt patterns and reliability.
 - Compare local vs hosted models (Ollama vs watsonx).
- **Day 2:** RAG
 - Local RAG notebooks (Ollama + Chroma).
 - RAG notebooks in labs-src (Elasticsearch, Chroma, Granite).
 - Accelerator ingestion + retriever + pipeline skeleton.
- **Day 3:** Orchestration & agents
 - Agent notebook calling the accelerator API.
 - Accelerator API & (optional) Streamlit UI.
 - Governance & Orchestrate notebooks.

PART 2 – GROUP REFLECTION

Suggested questions for the group:

- What was most surprising across the three days?
- Biggest “aha” moments?
- Hardest parts of the labs?
- Where did you hit friction with:
 - Environments?
 - Libraries / SDKs?
 - Conceptual pieces (RAG, agents, governance)?

Capture answers on a shared board or document.

PART 3 – OPEN Q&A & TROUBLESHOOTING

Use this time to address remaining technical issues:

- Ollama / watsonx environments:
 - Model loading, API keys, network.
- Accelerator service:
 - /ask endpoint behaviour.
 - Vector DB connectivity.
 - Logging & configuration.
- Agent notebooks:
 - Tool selection.
 - JSON formatting issues.
 - Error handling.

Encourage participants to share **their code** and walk through solutions live.

PART 4 – FROM WORKSHOP TO REAL PROJECTS

Discuss how to turn the accelerator into something “real”:

- **Production microservice**
 - Hardened FastAPI app.
 - Config via `service/deps.py` and environment variables.
 - Health checks, observability, logging.
- **Scalable ingestion pipeline**
 - Batch extraction / chunking / embedding.
 - Scheduling (nightly or event-driven).
 - Monitoring index size and freshness.
- **Integration**
 - CI/CD pipelines using `Makefile`, `pyproject.toml`, `Dockerfiles`.
 - Deployment targets: Kubernetes / OpenShift / Cloud Foundry / VM.
 - Governance and evaluation workflows with `watsonx.governance`.

PART 5 – SUGGESTED NEXT STEPS

Potential directions for deeper dives:

- **Retrieval & ranking**
 - Hybrid search (lexical + semantic).
 - Rerankers and scoring strategies in `retriever.py`.
- **Prompting & safety**
 - Advanced prompt strategies in `prompt.py`.
 - Safety / refusal patterns and guardrails.
- **Multi-tenancy**
 - Tenant-aware configuration in `service/deps.py`.
 - Per-tenant indices and model settings.
- **UI & UX**
 - Richer Streamlit UI (`ui/app.py`).
 - Conversation history, feedback buttons, source highlighting.

Provide pointers to:

- Internal docs and repos.
- IBM public samples (`labs-src`, GitHub repos).
- Governance and Orchestrate documentation.

OPTIONAL: FEEDBACK FORM

Share a link or QR code to a short survey covering:

- Overall satisfaction.
- Clarity of explanations.
- Lab difficulty.
- Which topics participants want more of.

Use this to plan the **capstone day** and any follow-up sessions.