

3.0 AGENTIC AI & ORCHESTRATION OVERVIEW

LEARNING OBJECTIVES

By the end of Day 3 (morning theory block), you should be able to:

- Explain what **Agentic AI** is and how it differs from “single-shot” LLM calls.
- Describe common **agent patterns**: tool calling, ReAct, plan–act, multi-agent and graph-based flows.
- Map these patterns onto concrete frameworks:
 - **CrewAI** – multi-agent orchestration in Python.
 - **Langflow** – visual builder for LangChain-style flows.
 - **LangGraph** – graph-based orchestration for complex workflows.
 - **watsonx Orchestrate** – IBM’s production-grade agent platform.
- Understand how your **RAG accelerator** turns into a reusable tool for agents.

3.1 WHAT IS AGENTIC AI?

So far we mostly used LLMs like functions:

Prompt in → answer out

Agentic AI adds **reasoning + action**:

- The LLM doesn't just answer — it:
 - Decides **which tools to use** (APIs, RAG services, calculators...).
 - Plans a sequence of **steps**.
 - Executes tools, reads results, and continues reasoning.
- The result is an **agent**:
 - Has a goal (“help user with workshop questions”).
 - Has **capabilities** (tools, collaborators, knowledge bases).
 - Uses an LLM for planning and reflection.

Agentic AI is powerful when:

- You need **multi-step workflows** (e.g. search → filter → call RAG → summarize).
- You need to integrate with **existing systems** (ticketing, CRM, data lakes).
- You want **traceability**: which tools were used, which docs were read, etc.

3.2 CORE AGENT PATTERNS

We'll refer to these patterns throughout Day 3:

1. TOOL-CALLING AGENT

- LLM chooses **one tool** at a time based on user input.
- Example:
 - `rag_service_tool` (question) → calls your accelerator / ask endpoint.
 - `calculator_tool` (expression) → safe arithmetic evaluation.
- Agent picks a tool, calls it, and formats the result back to the user.

2. REACT (REASON + ACT)

- LLM alternates between **thinking** and **acting**:
 - **Thought:** “I should call the RAG service to get context.”
 - **Action:** Use tool `rag_service_tool`.
 - **Observation:** Tool output.
 - **Repeat...**

Many frameworks (CrewAI, LangGraph, watsonx Orchestrate “react style”) build on this idea.

3. PLAN-ACT

- LLM plans a sequence of steps, then executes them:

3.3 FRAMEWORK TOUR

In the morning we'll conceptually walk through four frameworks that implement these patterns.

3.3.1 CREWAI

- **What it is:** A Python library for multi-agent “crews”.
- **Mental model:**
 - Agent – role, goal, backstory, tools.
 - Task – what needs doing, expected output.
 - Crew – group of agents + tasks + process (Process. sequential, Process.hierarchical,...).
- **Where it shines:**
 - Faster prototyping of **multi-agent** patterns.
 - Narrative workflows (research → writing → editing).
- **Workshop angle:**
 - A single CrewAI agent using your **accelerator RAG API + calculator**.

3.3.2 LANGFLOW

- **What it is:** A visual builder for LangChain flows.
- **Mental model:**
 - Drag-and-drop components (LLM, retriever, tools, routers, prompts).
 - Connect them as a **graph**.

3.4 MORNING FLOW (APPROX. 4H)

Suggested agenda (adapt to your group):

1. Intro & recap (30–45 min)

- Why agents on top of RAG?
- Quick recap of RAG architecture from Day 2.

2. Agent patterns walkthrough (45 min)

- Tool-calling, ReAct, plan–act, multi-agent, graph-based.

3. Framework lightning demos (90–120 min)

- **CrewAI mini-demo:**

- One agent with a “support engineer” role.
 - Tools: calculator + stubbed RAG tool.

- **Langflow mini-demo:**

- Visual RAG chain that calls an LLM and shows citations.

- **LangGraph mini-demo:**

- Simple graph: retrieve → generate.

- **Orchestrate concept demo:**

- Show YAML / ADK structure for a “Hello World Agent”.

4. Bridge to afternoon labs (15–30 min)

- Show how the accelerator / ask endpoint becomes the key tool.

3.5 HOW THIS CONNECTS TO LABS

- Day 2 → Day 3 bridge:
 - Day 2 gave you a **production-like RAG service** (/ask endpoint).
 - Day 3 gives you **agents** that call that service as a tool.
- Lab 3.1 – Local Agent in simple-watsonx-environment:
 - Agent has two tools: RAG service, calculator.
 - Planner LLM chooses which tool to call, then composes final answer.
- Agent frameworks (**CrewAI / LangGraph / Orchestrate**):
 - You can re-express the same logic in different frameworks:
 - CrewAI for Python multi-agent setups.
 - LangGraph for stateful workflows and evaluation.
 - Orchestrate for production deployment and governance.

By the end of Day 3, your mental model should be:

Docs → RAG (Day 2) → Agent on top (Day 3) → Orchestrated & governed in watsonx (beyond the workshop).

3.2 BRIDGE TO WATSONX ORCHESTRATE & GOVERNANCE

LEARNING OBJECTIVES

- Understand how notebook patterns translate to watsonx Orchestrate & governance tools.
- Recognize where policies and monitoring fit.

MAPPING ACCELERATOR + AGENT TO ORCHESTRATE

Think of watsonx Orchestrate as a **platform** that can host the patterns you built in notebooks.

- Conceptual mapping:
 - accelerator/service/api.py → a reusable “RAG answer” action (tool / service).
 - Agent notebook → blueprint for a multi-step orchestrated workflow.

EXAMPLE ORCHESTRATE FLOW

1. Receive a user request (chat, API, UI).
2. Decide whether to:
 - Call the **RAG service** (accelerator / ask endpoint).
 - Call a calculator or other utility tool.
 - Escalate to a **human** (handoff / ticket).
3. Log all steps and outcomes for governance.

In Orchestrate terms:

- **Agent** = top-level orchestrator that plans and delegates.
- **Tools** = your APIs (RAG, calculator, external systems).
- **Connections** = how tools authenticate to external services.
- **Knowledge bases** = RAG-ready corpora defined in Orchestrate.

GOVERNANCE & EVALUATION

watsonx.governance adds a layer of **control and insight** over your agents and models:

- **Model & asset catalog**
 - Track which models, prompts, tools, and agents exist.
 - Versioning and metadata.
- **Policies**
 - Allowed models / endpoints.
 - Risk profiles and approval workflows.
 - For example:
 - “Customer-facing agents may only use models with a given risk rating.”
- **Evaluation Studio**
 - Run experiments on your agents / RAG flows.
 - Compare LLMs, prompts, retrieval strategies.
 - Track metrics like faithfulness, answer relevance, content safety.
- **Monitoring**
 - Track runs and metrics over time.
 - Detect drift or degradation.
 - Investigate failures and edge-cases.

USING EXISTING NOTEBOOKS AS A BRIDGE

You already have governance-related notebooks in labs-src and accelerator:

- labs-src/ibm-watsonx-governance-governed-agentic-catalog.ipynb:
 - Shows how to:
 - Register models & tools in a governed catalog.
 - Create and use governed tools (e.g. PII detectors, jailbreak detectors).
- labs-src/ibm-watsonx-governance-evaluation-studio-getting-started.ipynb:
 - Shows how to:
 - Define evaluation datasets.
 - Compute metrics like context relevance, faithfulness, answer relevance.
- accelerator/assets/notebook/notebook>Analyze_Log_and_Feedback.ipynb:
 - Concrete example of:
 - Pulling logs from a deployed service.
 - Analyzing quality and user feedback.
 - Bridging logs → evaluation datasets.

Your Day 3 agent logs can be fed into these notebooks as a **first step** towards full governance.

FROM LOGS TO GOVERNANCE

Where do logs come from?

- `rag/pipeline.py`:
 - Can log retrievals and model calls (question, chunks, model, latency).
- `service/api.py`:
 - Can log user requests, status codes, errors.
- Agent notebook (Lab 3.1):
 - Can log tool choices, planner outputs, tool responses, final answers.

These logs can be:

1. Exported periodically as CSV/JSON.
2. Loaded into:
 - Evaluation Studio datasets.
 - Analysis notebooks (e.g. `Analyze_Log_and_Feedback.ipynb`).

Governance workflows can then:

- Detect drift in answer quality.
- Enforce thresholds:
 - “If faithfulness < 0.8, flag for review.”
- Support audit trails:

EXAMPLE USE CASES

Some concrete patterns where orchestrate + governance shine:

- **Governed HR assistant**
 - Uses RAG over HR policies.
 - Tools for detecting PII and harmful content.
 - Governed model + tool catalog and metrics.
- **Customer support bot with monitored outputs**
 - Agent routes between FAQ RAG, ticket creation, and human handoff.
 - Governance monitors answer relevance and safety.
- **Internal knowledge bot with evaluation cycles**
 - Weekly evaluation runs on a curated question set.
 - Results feed into prompt/model/index improvements.

HOW TO GO FROM LAB TO PRODUCTION

From Day 3 notebooks to real systems:

- Extract the best patterns from:
 - `agent_watsonx.ipynb` (agent loop + tools).
 - `accelerator service` (RAG microservice).
 - Governance notebooks (evaluation + catalog).
- Embed them into your delivery pipeline:
 - Use `Makefile` / `pyproject.toml` / Dockerfiles to package the accelerator.
 - Deploy accelerator API behind a stable endpoint.
 - Register it as a tool / connection in Orchestrate.
- Connect Orchestrate deployments to governance:
 - Ensure models and agents are registered in the catalog.
 - Route logs to Evaluation Studio or custom analysis.

DISCUSSION PROMPTS

To close the session, consider:

- Where would you use **orchestration** and **governance** in your organization?
- Which components from this workshop are closest to your production needs?
- What would you need to:
 - Onboard your own data?
 - Apply your internal policies?
 - Connect to your existing IT systems?

Capture these ideas — they are great seeds for the **capstone day** and follow-up projects.

3.3 RECAP & NEXT STEPS (2H DISCUSSION)

SESSION OVERVIEW

This is a 2-hour interactive recap and discussion to consolidate learning and plan next steps.

PART 1 – TIMELINE RECAP

Walk through the workshop story:

- **Day 0:** Environment setup
 - Two env repos (Ollama + watsonx).
 - accelerator repo.
 - labs-src reference notebooks.
- **Day 1:** LLM basics & prompting
 - Prompt patterns and reliability.
 - Compare local vs hosted models (Ollama vs watsonx).
- **Day 2:** RAG
 - Local RAG notebooks (Ollama + Chroma).
 - RAG notebooks in labs-src (Elasticsearch, Chroma, Granite).
 - Accelerator ingestion + retriever + pipeline skeleton.
- **Day 3:** Orchestration & agents
 - Agent notebook calling the accelerator API.
 - Accelerator API & (optional) Streamlit UI.
 - Governance & Orchestrate notebooks.

PART 2 – GROUP REFLECTION

Suggested questions for the group:

- What was most surprising across the three days?
- Biggest “aha” moments?
- Hardest parts of the labs?
- Where did you hit friction with:
 - Environments?
 - Libraries / SDKs?
 - Conceptual pieces (RAG, agents, governance)?

Capture answers on a shared board or document.

PART 3 – OPEN Q&A & TROUBLESHOOTING

Use this time to address remaining technical issues:

- Ollama / watsonx environments:
 - Model loading, API keys, network.
- Accelerator service:
 - /ask endpoint behaviour.
 - Vector DB connectivity.
 - Logging & configuration.
- Agent notebooks:
 - Tool selection.
 - JSON formatting issues.
 - Error handling.

Encourage participants to share **their code** and walk through solutions live.

PART 4 – FROM WORKSHOP TO REAL PROJECTS

Discuss how to turn the accelerator into something “real”:

- **Production microservice**
 - Hardened FastAPI app.
 - Config via `service/deps.py` and environment variables.
 - Health checks, observability, logging.
- **Scalable ingestion pipeline**
 - Batch extraction / chunking / embedding.
 - Scheduling (nightly or event-driven).
 - Monitoring index size and freshness.
- **Integration**
 - CI/CD pipelines using `Makefile`, `pyproject.toml`, `Dockerfiles`.
 - Deployment targets: Kubernetes / OpenShift / Cloud Foundry / VM.
 - Governance and evaluation workflows with `watsonx.governance`.

PART 5 – SUGGESTED NEXT STEPS

Potential directions for deeper dives:

- **Retrieval & ranking**
 - Hybrid search (lexical + semantic).
 - Rerankers and scoring strategies in `retriever.py`.
- **Prompting & safety**
 - Advanced prompt strategies in `prompt.py`.
 - Safety / refusal patterns and guardrails.
- **Multi-tenancy**
 - Tenant-aware configuration in `service/deps.py`.
 - Per-tenant indices and model settings.
- **UI & UX**
 - Richer Streamlit UI (`ui/app.py`).
 - Conversation history, feedback buttons, source highlighting.

Provide pointers to:

- Internal docs and repos.
- IBM public samples (`labs-src`, GitHub repos).
- Governance and Orchestrate documentation.

OPTIONAL: FEEDBACK FORM

Share a link or QR code to a short survey covering:

- Overall satisfaction.
- Clarity of explanations.
- Lab difficulty.
- Which topics participants want more of.

Use this to plan the **capstone day** and any follow-up sessions.